

HUMANITY

A framework for social organizations to build emergency apps that inform without shutting down our sensitivity.

Proposal fact sheet

Challenge Track	AI and MIL (the proposal also addresses the Community Impact track, by working together with civil defense, NGOs, and grassroots community organizations on the ground).
Proposal format	Apps/Websites in a hybrid format: a framework of principles and reusable technical components (equivalent to an educational toolkit for organizations) that takes shape in a reference application, HumanityJunín.
Hackathon theme	<i>Play Your Part: Youth Designing the Future of Media and Information Literacy (MIL)</i>

I. Executive summary

Problem. In emergencies, AI systems optimized for engagement display a real tragedy and a viral dance with the same "visual equality," while hallucinations and deepfakes threaten to manipulate humanitarian aid. The result is a young audience that consumes disaster content without the media and information literacy (MIL) skills needed to verify sources, distinguish spectacle from genuine emotion, and decide with their own judgment whether to share or donate.

Solution. Humanity is a framework of six ethical design principles, not a single app, that any social organization can follow to build its own emergency applications. It translates MIL competencies into concrete features: a reflective pause before sharing, image provenance verification (C2PA) to detect AI-manipulated content, questions about who is speaking and from where before trusting a source, and blockchain traceability of donations.

Impact and validation. The framework is validated with a real case, HumanityJunín, built after the earthquake of July 18, 2026, with bilingual Spanish-Quechua protocols and WCAG 2.1/2.2 accessibility. Its KPIs measure not only usage but media literacy learning: the rate at which the reflective pause is activated, AI-manipulated content detected and confirmed by the community, and "ethical recurrence" (people who, after using the app, go on to question content from other emergencies outside it).



II. About us

We are a complementary team: we bring together a technical, philosophical, and social perspective so that our proposal is not just a tool but a way of understanding people.

Flor Betzabeth Ampa Flores

Systems Engineering student and Philosophy student.

Nationality: Peruvian

Age: 23 years old

Why am I here? I believe technology has the power to make us more human; we need to redirect the meaning of making technology, because it is already present in every tiny aspect of our lives.

Tais Ampa Flores

Social Work student.

Nationality: Peruvian

Age: 22 years old

Why am I here? I want to contribute a perspective of community intervention, an ethics of care, and direct work with people affected by emergencies.

We are sisters, and we propose this solution, this framework, from our own experience. We grew up in a difficult environment surrounded by poverty and violence, and even so we **COULD NOT** ask for help, **WHY?** Because **WE WERE ASHAMED**, it wasn't pride — we didn't even know the meaning of that word; it was something we now recognize as a lack of sensitivity from those around us.

III. Problem statement:

How did the idea come about?

In an aesthetics class at the philosophy school, on an assignment given to us by Dr. Ana María, "*La razón estética: una propuesta para tiempos difíciles*" [*Aesthetic Reason: A Proposal for Difficult Times*], which discusses the "celebritization" of media, what counts as important on social media, and the need for an education in sensibility to address this problem. Although the author ultimately gave up on her own solution, I kept thinking about what could be done.

I came to the conclusion that, no matter how much we educate people, if the system they operate within leads them toward desensitization to what matters most — that anesthesia against pain that social media offers — real change will never happen. That's when I saw it from another angle: what truly needs to change is the way we create technology, especially with artificial intelligence, which is already part of our daily lives.

Problem

We live surrounded by screens that display everything in the same format: an earthquake, a cooking recipe, and a viral dance appear one after another, the same size, for the same duration. That "visual equality" comes at a cost: little by little, we stop feeling moved by real tragedies, because we see them as just another video within the spectacle.

Main causes:

- AI systems are designed for extreme personalization and amplification, prioritizing content that generates "engagement" (views and clicks) regardless of whether it's a real tragedy or a viral dance. This creates a bias where AI treats all information as merchandise

- Technology is regarded merely as a tool, when in reality it influences our lives far more than that.
- Consuming content produces satisfaction in us.
- AI hallucinations and deepfakes in disaster contexts. In an emergency, a fake AI-generated image can divert aid or dishonestly manipulate people's sensitivity (automated "poverty porn")

Main effects:

- Automatic responses, without pausing to actually think or feel what we're seeing. The result is serious: people know something terrible has happened, but they don't act, because they're already used to the idea that "everything is content."
- AI is seen as a mode of optimization, a more efficient way of doing things, but one oriented more toward economic and commercial-production motives.
- Faced with any kind of information, it becomes genuinely difficult to tell a genuine emotion apart from a reaction to spectacle.
- People are taught to stay informed without realizing what that process is actually doing to them.




Figure 1. From today's algorithmic model to Humanity's cyclical, ethical model.

Today, the technology we consume works as a closed, linear cycle: content — whether a tragedy or a viral dance — feeds an algorithm that optimizes “engagement”; the algorithm pushes that content to the user; the user reacts automatically; and that reaction becomes interaction data that feeds back into the same algorithm, reinforcing the pattern without any ethical pause along the way. Humanity proposes breaking that closed cycle and replacing it with an open, two-way one: development doesn't end with the user, but rather engages in dialogue with them through co-design and real feedback; that feedback always passes through an ethical review with human oversight (human-in-the-loop) before becoming an improvement; and those improvements return to development, not to an algorithm that only measures clicks. The core difference isn't technical but about how the cycle itself is designed: one optimizes for attention, the other for judgment.

IV. Objectives

General objective:

Design **Humanity**, a *framework* focused on creating applications that empathize with, respect, and do not manipulate their users — whether those users need to make a situation of vulnerability visible or wish to offer their support.

Specific objectives:

- Help organizations that want to support their community build more empathetic applications, more oriented toward the user as a human being, taking their dignity and needs into account.
- Create a model that can be used to build emergency applications for earthquakes, natural disasters, or fundraising campaigns for social causes.
- Inform and connect people without shutting down their sensitivity, denying their dignity, or manipulating their willingness to help or ask for help.
- Encourage a reflective pause before sharing sensitive content, moving from automatic reaction to a conscious decision.
- Help users distinguish between a genuine emotion and a "spectacle" reaction to disaster content.
- Offer an accessible, inclusive experience for children, young people, adults, older adults, and people with disabilities.
- Make a replicable framework — not just a closed application — available to other organizations, so the impact multiplies.
- Include guides within the framework so that apps can help young people detect when disaster content has been generated or manipulated by AI.

V. Target audience

Direct users: young people and adults who consume content about emergencies on social media and applications, and who want to stay informed and help responsibly.

Partner organizations: NGOs, grassroots community organizations, civil defense, and youth groups that need a reliable tool for communicating emergencies and raising funds.

Groups with specific needs: children, older adults, and people with visual or motor disabilities, for whom accessibility adjustments are built into the design from the start.

By thinking of it as a framework (not just an app), **Humanity** can adapt to different contexts: an organization responding to an earthquake — Junín (Peru); another raising funds for a community affected by, for example, the cold wave ("friaaje") in Peru's southern highlands; or a youth group wanting to launch a social campaign, such as a Christmas "chocolatada" (hot chocolate drive).

VI. Concept and prototype: The Humanity framework

Humanity is not a single application, but rather a "toolbox" (a framework) built around six principles that any organization can follow to build its own aid app, so that communication and information not only fulfill the role they exist for, but also educate everyone involved in sensitivity. This way, instead of each team starting from zero, everyone shares the same ethical and design foundation.

Principles:

- a) Development must be oriented toward respect and dignity for all users.
- b) It must avoid shutting down or manipulating the sensitivity of users and developers.

- c) Use technologies that promote transparency.
- d) Avoid deceptive design.
- e) Value questions within the system.
- f) Aim to shape people, not just capture users.

How does it align with the proposed hackathon?

The hackathon calls for proposals that apply artificial intelligence ethically and responsibly to real social problems, in line with UNESCO's Recommendation on the Ethics of Artificial Intelligence (2021). «Humanity» does not propose using AI as an extra technological attraction, but rather makes it the very core of the problem it seeks to solve: it is AI-driven recommendation systems that today treat a tragedy and a viral dance with the same “visual equality,” and it is AI-generated hallucinations and deepfakes that threaten to manipulate humanitarian aid in disaster contexts. That's why the framework responds to the pillars UNESCO itself requires for ethical AI: human oversight (human-in-the-loop) in detecting “spectacle” content, transparency through provenance verification (C2PA) and public donation records, non-maleficence by avoiding deceptive design and emotional exploitation, and proportionality by leaving the final ethical decision in people's hands rather than the algorithm's. Humanity translates those general principles into a concrete checkpoint that any team at this hackathon, or any organization afterward, can apply to its own socially impactful AI project — directly answering the replicable spirit the call for proposals seeks.

Beyond AI ethics, Humanity is anchored directly in UNESCO's Media and Information Literacy (MIL) framework, the central theme of this hackathon ("Play Your Part: Youth Designing the Future of Media and Information Literacy"). Each feature of the framework corresponds to a recognized MIL competency: the **critical evaluation of sources** (asking who is speaking, from where, and with what intention, before trusting a piece of content); the **verification of the provenance and authenticity of information** (C2PA standards to distinguish AI-generated or manipulated content from a real record); and **critical thinking about the media** (the reflective pause that interrupts the automatic reaction to "spectacle" content). Humanity doesn't just want its users to "consume" emergency information; it wants them to develop and carry these MIL competencies beyond the app itself, as reflected in the specific objective of "shaping people, not just capturing users," and as measured in the MIL competency KPIs detailed further below.

Key aspects you can include in your APP built with Humanity

Real emotions, not spectacle	The app uses an AI model that helps the person notice how much what they saw that day actually mattered to them, instead of being swept along purely by visual impact. It doesn't push the person to act immediately: it first seeks to help them understand, and it respects each person's own decision about whether they want to help. It's designed for all ages and for people with disabilities (guided reading, button-based navigation, visual accessibility).
Pain isn't "flattened"	Each emergency and each affected person is presented differently, avoiding the generic short-video format that makes a flood look the same as a viral dance. First-person testimonies are prioritized: who needs help, exactly what they need, and how they can be supported.

Transparency and critical thinking	Before trusting a source, the app invites the user to ask who is speaking, from where, and with what intention (a decolonial approach). Donations are recorded using blockchain technology so that anyone can verify that the money reaches where it's supposed to.
No tricks, no "deceptive design"	The interface avoids dark patterns (buttons that pressure people to donate, fake urgency counters, etc.). Every action, however small — even tapping a button — is designed with care, because even minimal gestures shape values.
A pause before reacting	Before sharing content, the app asks a brief question ("Why do I want to share this?", "Have I already verified it's real?"). That small bit of friction helps shift an automatic reaction into a more conscious decision.
Shaping people, not just users	The ultimate goal isn't just getting clicks or donations, but helping each person build their own ethical judgment: when to share an image, how to act in the first few minutes after seeing intense content, and how to become an informed, sensitive "MIL hero."

About the Framework

Before development

- Build multidisciplinary teams that analyze their target audience. A psychologist knows more about depression than a designer, and a designer knows more about technical matters than a psychologist.
- Hold integration events that help build contact with end users. Maintain a perspective of otherness.
- Analyze the country's political context. For example, in Peru people trust social-media news influencers more than the state's own official channel.
- Gather all the possible variables that could influence the development of this application. Search not only official sources but also social media, and use the “human in the loop” method to rule out videos that may be AI-manipulated or taken out of context.

During development

- Ethical checkpoints at every sprint: before accepting a feature, the team checks it against Humanity's six principles (Does it respect dignity? Does it avoid deceptive design? Can it be validated transparently?).
- Co-design with partner organizations and with people from the affected groups, not just internal testing: those who will ask for or give help take part from the earliest prototypes.
- Open documentation of the code and design decisions, so another organization can audit or adapt the framework without depending on the original team.

After development

- Impact metrics centered on wellbeing, not just usage: for example, how many people took the reflective pause before sharing, rather than measuring only clicks or screen time.
- Permanent feedback spaces with user communities, so the app keeps learning from real experience and not only from usage data.
- Periodic ethical audits of the framework, reviewing whether new features (such as AI integrations) still respect the six principles.

Roles and responsibilities

The project lead, together with whoever holds the ethical perspective (psychology, social work, or philosophy), decides which feature is approved after each sprint's ethical checkpoint; the development team implements it following the defined technical standards (accessibility, provenance verification, encryption); and the partner organizations, together with end users, take part as validators through co-design, with a real voice to stop a feature before launch if it fails to respect people's dignity.


Communication

Channels and forms of interaction within the group: brief follow-up meetings between the technical and social teams, an open decision log so that another organization adapting the framework understands the reasoning behind each choice, and periodic spaces for direct feedback with user communities, so communication isn't only internal but also reaches those who use the app.

Processes

- **Rules and guidelines:** Official steps for carrying out the day-to-day work: before accepting any feature, the ethical checkpoint is applied against Humanity's six principles; every new feature goes through co-design with the partner organization and, whenever possible, with people from the affected groups; and every design decision is documented openly, so the day-to-day work is auditable and replicable by another team without depending on the original one.
- **Workflow:**

The logical order from when an idea begins to when it's delivered:



Continuous improvement: Reviews to fix errors and adjust the system: periodic ethical audits that verify whether new features (such as AI integrations) still respect the six principles, impact metrics centered on wellbeing and not only usage, and permanent feedback spaces with communities so the framework keeps learning from real experience.

Technology

Base tools: Reusable software, libraries, or code.

- **We propose an open, replicable stack:** interfaces in React or Flutter following WCAG 2.1 AA accessibility standards (guided reading, high contrast, button-based navigation); a modular backend in Python (Django/FastAPI) or Node.js.

Detecting "spectacle" content versus genuine emotion relies on natural language processing and sentiment analysis models, always with human oversight (human-in-the-loop) to prevent AI from deciding on its own what is or isn't worthy of attention.

To curb deepfakes in emergencies, we integrate image provenance verification using content credential standards (C2PA / Content Credentials).

Infrastructure: Servers and supporting networks.

- Cloud hosting with an open, federated data architecture, so that small organizations (a regional civil defense unit, a youth group) aren't dependent on a single closed platform. Donations are recorded on a low-cost public blockchain network, so that anyone can track that the money reaches where it's supposed to go. Sensitive data about affected people — especially children and people with disabilities — is encrypted end-to-end and stored under a minimization principle: only what is strictly necessary to provide aid is kept.

Technical standards:

- Mandatory security and quality protocols, including: WCAG 2.1/2.2 for web and mobile accessibility; "privacy by design" principles aligned with Peru's Personal Data Protection Law (Law No. 29733) and equivalent international frameworks; and, as the normative foundation of the entire framework, UNESCO's Recommendation on the Ethics of Artificial Intelligence (2021), which guides the criteria for transparency, human oversight, and non-maleficence in every AI component that Humanity uses.

VII. What would an application built with this framework look like?

*Let's imagine "**HumanityJunín**", an application built by a regional civil defense unit after an earthquake in Junín that took place on July 18, 2026, following the framework step by step.*

Before development

What is Junín?

Junín is a department located in the central zone of the Peruvian Andes, in the country's central highlands. Its 44,197 km² territory covers both the Andean region to the west and the Amazon rainforest to the east, and it borders Pasco to the north, Lima to the west, Ayacucho and Huancavelica to the south, Cusco to the southeast, and Ucayali to the northeast. Its capital and most populous city is Huancayo, located in the Mantaro Valley.

The official language is Spanish, alongside Quechua (the Huanca and Yaru variants), and other languages such as Asháninka, Caquinte, Nomatsiguenga, and Yanesha, with Quechua being the predominant native language among much of the population.

The climate in the highlands is cold and dry, with sharp temperature swings between day and night; in the higher areas, nighttime temperatures can drop to around 5°C, while daytime temperatures reach close to 25°C — a fact that worsens the situation for those who lost their homes in the earthquake, many of whom are left exposed to intense cold. [BloggerViajaraperu](#)

Internet access in these rural, high-Andean areas is limited, which is a central challenge for the app's design.

Material damage and response measures recognized so far

Reporte de daños en los distritos mas afectados por el sismo

12/04/17

Departamento de Junín		Vida y Salud				Daños Materiales		
Provincia ♦	Distrito ♦	Damnificados ♦	Afectados ♦	Fallecidos ♦	Heridos ♦	Destruída ♦	Inhabitable ♦	Afectada ♦
Chupaca	Chongos Bajo	2077	1792	5	21	129	633	587
	San Juan de Yscos	190	450			1	112	313
	Ahuac	54	123				30	78
Huancayo	Chongos Alto		181					101
	Chupuro	1511	447		130	57	513	171
	Huayucachi	593	381			165	13	151
	Viques	591	381			72	254	83
	Huacrapuquio	259	581	1		1	117	243

Additional cultural and climate context

Communities in the area mostly wear traditional "polleras" (layered skirts), a culturally relevant element for designing culturally appropriate content and interfaces.

In addition, the region faces a water deficit linked to the El Niño phenomenon, adding another layer of vulnerability on top of the one already caused by the earthquake.

The State's response to the situation

Distribution of humanitarian aid

The resources sent are managed in coordination with the three levels of government and include both public and private assistance:

- **INDECI shipment:** Peru's National Civil Defense Institute (INDECI) arranged for the shipment of 111.04 tons of humanitarian aid goods to the Junín region, managed with support from the World Food Programme. The shipment consists of 90.94 tons of non-food goods and 20.11 tons of food goods. This shipment came on top of an earlier one of 7.74 tons of food aid intended for 1,250 affected people. [El Peruano + 2](#)
- **Ministry support:** agencies such as the Ministry of the Environment and SERNANP are distributing supplies directly in affected areas such as Chupuro, while the Ministry of Women and Vulnerable Populations and the Ministry of the Interior, together with the National Police, have mobilized humanitarian flights and ground transport carrying food, clothing, and basic necessities. [Revista Signature](#)
- **Military and emergency deployment:** the Peruvian Army deployed to support search and rescue, debris removal, public safety, and humanitarian assistance, while health authorities set up three temporary medical posts in Pumpuya, and the Junín Regional Government opened a temporary shelter with blankets, tents, beds, and mattresses. [Wikipedia + 2](#)
- **Coordinated international aid:** INDECI, together with the Junín Regional Government and the municipalities of Chongos Bajo and Chupuro, organized the distribution of goods donated by UNHCR, including thermal blankets, sleeping mats, solar lights, jerrycans, buckets, and prefabricated units. [Government of PeruGovernment of Peru](#)

- **Private aid channels:** organizations such as ADRA Peru continue running collection drives and distributing hygiene and warmth kits in districts like Chongos Bajo, Viques, and Los Angeles.

Additional consideration

The app's design must account for the ongoing possibility of another earthquake in the country; in fact, the Chupaca area already recorded a minor aftershock in the weeks following the main event, reinforcing the need for "HumanityJunín" to also work as a prevention and early-warning tool, not only a response one.

User research questions (human-centered design)

Who would support this cause?

There isn't a single source that measures this precisely, so instead of claiming "most of the national population," it's better to segment potential supporters:

- Urban populations from other regions (especially Lima), with the financial means to donate.
- The Peruvian diaspora community abroad.
- NGOs and religious organizations already active in the area (ADRA Peru, Cáritas, TECHO, CARE Peru).
- Private companies with corporate social responsibility programs.
- Other Andean communities that identify culturally with Junín.

What is a person looking for when helping in the aftermath of an earthquake?

They're looking to translate an emotion (helplessness, empathy, guilt over the distance) into a concrete, verifiable action. They want confirmation that their help reaches a real person, not an opaque system.

What do they want to feel?

They want to feel that their contribution had a tangible effect, that it wasn't in vain; they're looking for emotional closure (knowing "what it was used for") and a sense of human connection, not just a transaction.

What can they learn about sensitivity?

They can learn not to impose their own urban/Western logic on a community with a different language (Quechua), different priorities (warmth before "generic kits"), and a different relationship with time and extreme weather. They can also learn to avoid paternalistic charity and to recognize the agency of the affected community.

What does someone from Peru's south-central highlands need to learn about an earthquake?

They need clear, culturally adapted (bilingual Spanish-Quechua) protocols on: what to do during the earthquake, how to put together an emergency kit, where to find safe zones given the Andean terrain (landslides, adobe housing), and how to act during aftershocks, given that the region already has a history of seismic activity from the active fault of the Huaytapallana mountain range.

What does an average Peruvian tend to assume about these seismic events?

They tend to assume "the worst has already passed" after the first quake (underestimating aftershocks), that state aid will arrive quickly and in an orderly way, and that adobe or makeshift homes will hold up as they have before. They also tend to normalize seismic risk because of how frequently minor earthquakes occur in the country.

What information is needed from the person offering support?

For the app to be useful and responsible: location (to link donations to specific areas), type of support they can offer (money, goods, volunteering, spreading the word), availability frequency, and consent to receive updates on how their contribution was used (traceability).

Verifiable context questions

What is the country's economic situation?

Peru's economy grew 3.5% in the first quarter of 2026, driven by private spending, and BBVA Research projects growth of 3.1% for the year. The BCRP (Peru's central bank) projects growth of close to 3%, sustained by historically high copper and gold prices, although cumulative inflation from January to April 2026 reached 3.53%, exceeding the 1%–3% target range set by the Central Reserve Bank. The exchange rate is expected to close the year between S/3.20 and S/3.30 per dollar. Structurally, Peru reduced poverty from 60% in 2002 to 24% in 2013, and its GDP per capita exceeded US\$8,400 in 2024, though significant regional inequality persists — precisely the kind of gap that areas like Junín reflect most starkly after a disaster. [BBVA Research + 4](#)

What is the political context?

Keiko Fujimori became President of Peru on July 28, 2026, after winning the general election by a very narrow margin over Roberto Sánchez, becoming the first woman elected president of the country by direct vote. She takes office after a decade marked by successive instability, with nine different people having held the presidency during that period. Especially relevant to this project, in her inaugural address Fujimori named mitigating the El Niño phenomenon and public safety as national emergency priorities, placing the response to climate and seismic disasters at the center of the new government's agenda — a key fact for thinking about how "HumanityJunín" could (or could not) align with current public policy. [Wikipedia + 3](#)

During development

Technology and design based on Humanity's principles

"HumanityJunín" reuses the open stack defined in the framework and adapts it to the constraints identified in the diagnosis: a React Native or Flutter frontend (to cover low-end Android devices, the most widely used operating system in the central highlands) with offline-first support through local storage and deferred synchronization while connectivity is limited; a modular backend in Python (FastAPI) for content-verification services and in Node.js for real-time messaging; NLP and sentiment-analysis models with human oversight to distinguish "spectacle" content from genuine testimonies; image provenance verification using C2PA/Content Credentials standards to curb deepfakes about the earthquake; and donation records on a low-cost public blockchain network.

Every decision on this list first goes through the ethical checkpoint: for example, facial recognition to identify affected people was ruled out because it would violate the principles of dignity and privacy by design.





You can visit the link to the simulated prototype:

<https://sites.google.com/unsu.edu.pe/humanityjunin?usp=sharing>

App features according to Humanity's six principles

- **Respect and dignity:** profiles of affected families written together with them (co-design), never as generic “cases”; explicit consent before publishing any photo or testimony.
- **Not shutting down sensitivity:** a limit on the number of testimonies shown per session and a variety of formats (text, audio, photo) to avoid visual habituation to the same short-video format.
- **Transparency:** blockchain traceability for every donation and a visible C2PA verification seal on every image.
- **Avoiding deceptive design:** no fake urgency counters or buttons pressuring people to donate; the suggested amount is never preselected by default.
- **Valuing questions:** the reflective pause and a bilingual FAQ space about what actually happened in Junín, written together with the community.
- **Shaping people, not just capturing users:** a brief “own judgment” module after every donation or share attempt, which never appears identically twice.

Phrases that reinforce the donor's sensitivity learning

These phrases don't appear as intrusive pop-ups, but are integrated into decision points: on the reflective-pause screen (before sharing), when confirming a donation, and on the emotional-closure screen (after donating). Examples:

- Before sharing: “What you're about to share is the real story of a family in Chupuro. Are you sharing it because you care, or because it's shocking?”
- When confirming the donation: “Your contribution doesn't end here: you'll be able to see, with a name and a date, what it was used for.”
- In the emotional closure: “You helped a real family, not an urgency counter. Next time you see a tragedy on your screen, you already know what to ask yourself.”

Design rationale based on the Junín analysis

Each decision above responds directly to the earlier diagnosis: bilingual Spanish-Quechua support and audio respond to the fact that Quechua is the predominant native language in the area; the offline-first design responds to the limited internet access identified in rural, high-Andean areas; the aftershock and safe-zone module responds to the fact that the active fault of the Huaytapallana mountain range makes a new event likely, as already confirmed by the minor aftershock in Chupaca; prioritizing warmth over “generic kits” responds to nighttime temperatures as low as 5°C recorded in the highlands; and blockchain traceability responds directly to what the user research identified as the core need of anyone donating: confirming that their help reaches a real person and not an opaque system.

In the end, whoever used “HumanityJunín” doesn't walk away with just the feeling of having clicked something: they walk away with their own judgment for the next time they see a tragedy on their screen (principle f). That same kit of components, principles, and technical standards can be reused by another organization for a flood campaign or to raise funds for a different social cause, without starting from scratch.

After development

KPIs for effective aid to the population of Junín

- **Amount and type of aid actually delivered** and verified on the blockchain, against the amount pledged (traceability compliance rate).
- **Average time between a donation and its arrival** to the beneficiary person or family.
- **Number of families in Chupuro, Chongos Bajo, Viques, and Pumpuya** reached directly through the app, compared with existing state and NGO channels.
- **Use of the prevention and early-warning module** (people who complete the bilingual Spanish-Quechua protocols) for aftershocks.

KPIs for genuine awareness-raising (not just usage)

- **Rate at which the reflective pause is activated** before sharing sensitive content, out of the total number of share attempts.
- **Change in decision time** between viewing content and donating or sharing (a slower, more conscious decision is interpreted as a positive signal, not a negative one, unlike conventional immediate-conversion metrics).
- **Brief post-use surveys** on whether the person felt they understood the situation in Junín, not just that they “clicked” to help (a metric of emotional closure and human connection).
- **Ethical recurrence:** people who, after using the app once, report having questioned content from other emergencies outside the app (transfer of the learned judgment, principle f).
- **Detection of AI-manipulated content** reported and confirmed thanks to provenance verification (C2PA), as an indicator of the user community's critical literacy.

MIL competency KPIs (media literacy learning)

Unlike the previous KPIs, which focus on trust and donation traceability, these measure whether users actually acquire and transfer MIL competencies, not just whether they use the app:

- **Accuracy in source verification:** percentage of users who, given a brief in-app micro-test (real content vs. AI-manipulated content), correctly identify the manipulation, measured before and after a period of use.
- **Transfer outside the app:** percentage of people who, in follow-up surveys, report having questioned or verified content about other emergencies outside Humanity (measures whether the MIL competency transfers to the broader media ecosystem and doesn't only occur within the app).
- **Self-perceived judgment:** change, on a brief pre/post-use scale, in how much the person trusts their own ability to distinguish spectacle from genuine emotion and to identify who is speaking and with what intent in a piece of content.
- **MIL competency coverage in the ethical checkpoint:** proportion of new features that, at each sprint's checkpoint, explicitly incorporate at least one MIL competency (source verification, bias assessment, distinguishing information from opinion), and not only dignity or transparency criteria.

Areas for improvement

- Strengthen the app's offline and low-bandwidth performance, given the limited internet access in Junín's rural, high-Andean areas identified in the initial diagnosis.
- Expand validation with people with visual or motor disabilities beyond internal testing, involving real users in every iteration.

- Evaluate the cost and environmental footprint of the chosen blockchain network, and migrate to lower-energy alternatives if transaction volume grows.
- Periodically review whether the ethical checkpoint remains sufficient as new AI capabilities are integrated, preventing the pace of technical development from outrunning ethical reflection.

VIII. Sustainability

Humanity sustains itself financially by combining three sources that don't depend on a single donor:

1. Partnerships with organizations already present in the area (civil defense, NGOs, regional governments), which contribute local knowledge and distribution channels.
2. A small, transparent administrative percentage on donations processed, visible within the app itself thanks to the blockchain record.
3. Licensing or technical support for organizations adopting the framework for their own contexts (a municipality raising funds for a cold wave, a youth group organizing a Christmas campaign), not charging for the framework itself — which remains open — but for implementation support.

IX. Scalability

Being a set of principles, reusable components, and technical standards rather than a single closed application, Humanity scales horizontally: every new organization that adopts it adds its own instance (“HumanityJunín,” “HumanityFriaaje,” “HumanityInundaciones”) without the original team needing to grow at the same pace. What does need to grow is the community of practice around the framework: documentation, ethical-checkpoint templates, and a library of accessible components (WCAG 2.1/2.2) that any team can reuse.

Minimum team required

- **1–2 full-stack developers** (React or Flutter + backend in Python/Node.js) per instance of the framework.
- **1 UX/UI designer** with training or support in accessibility and ethical design.
- **1 social sciences professional** (social work, psychology, or philosophy) to sustain the ethical checkpoint and community co-design.
- **1 partnerships/community lead** who liaises with the partner organization and local authorities.
- **Shared technical support** (security, C2PA provenance verification, blockchain infrastructure) that can be a centralized service used by several instances of the framework at once, instead of being duplicated within each organization.

Implementation roadmap for Peru and Latin America

- **Phase 1: Pilot (0–6 months).** Implement “HumanityJunín” with a real partner organization, measure the first wellbeing KPIs, and adjust the framework based on the experience.
- **Phase 2: National consolidation (6–18 months).** Adapt the framework to another type of Peruvian emergency (cold wave in the southern highlands, flooding in the rainforest) to test its replicability in contexts other than earthquakes.
- **Phase 3: Regional expansion (18–36 months).** Offer the framework, already documented and audited, to organizations in other Andean and Latin American countries facing similar challenges (multilingualism, limited connectivity, high seismic or climate exposure), always prioritizing local co-design over a literal translation of the Peruvian app.

X. The creativity

What's creative about «Humanity» isn't a single isolated feature, but the shift in perspective it starts from: instead of asking “how do we get more people to donate or share?”, the project asks “how do we get the person donating or sharing to do so while genuinely feeling and understanding what's

happening?” That shift in the question is what sets Humanity apart from just another fundraising product.

- **A framework, not an app.** Instead of competing with other emergency apps, Humanity positions itself as an ethical and technical foundation that any organization can inherit, multiplying impact without multiplying the effort of designing ethics from scratch.
- **Friction as design, not as a flaw.** Most apps try to eliminate all friction to maximize conversion; Humanity introduces a minimal, deliberate friction (the reflective pause) at exactly the point where conventional design would try to eliminate it.
- **AI as a mirror, not a judge.** The “spectacle vs. genuine emotion” detection model doesn't decide for the user or censor content: it returns a question, leaving the ethical decision in human hands (human-in-the-loop) — something uncommon in social-impact apps.
- **Born from lived experience, not from a brief.** The project starts from a personal story of shame and a lack of sensitivity from those around us, not from an abstract business problem; that allows us to design for the dignity of the person asking for help with the same care as for the person giving it.
- **A real case as a laboratory.** “HumanityJunín” isn't a generic mockup: it's built on a situated analysis (the Quechua language, adobe housing, limited connectivity, aftershocks), which forces the design's creativity to solve real constraints, not just aesthetic ones.

XI. The viability

«Humanity»'s viability rests on the fact that it doesn't propose building a closed platform from scratch, but rather an open framework that draws on already-proven technology, on organizations already operating in the field, and on existing international standards. That reduces the project's risk across three dimensions:

- **Technical viability.** The proposed stack (React/Flutter, a modular backend in Python or Node.js, C2PA provenance verification, donation records on a low-cost public blockchain network) uses open-source, well-documented tools with a reasonable learning curve for hackathon teams and for junior developers whom partner organizations might hire afterward.
- **Operational viability.** The framework doesn't depend on a single team: open documentation, ethical checkpoints, and defined technical standards allow another organization (a regional civil defense unit, a youth group, an NGO) to adopt Humanity without depending on whoever originally designed it.
- **Economic viability.** The initial cost is concentrated in low-cost cloud infrastructure and design/development hours; since it's a replicable framework rather than a one-off closed app, the marginal cost of each new implementation (“HumanityJunín,” “HumanityFriaaje,” etc.) decreases, because it reuses components, principles, and standards already built.
- **Legal and ethical viability.** By aligning from the design stage with UNESCO's Recommendation on the Ethics of Artificial Intelligence (2021) and with Peru's Personal Data Protection Law (Law No. 29733), Humanity doesn't need to rework its architecture to comply with regulations later: compliance is a design principle from the very first sprint.
- The greatest risk isn't technical but rather sustaining the ethical perspective over time; that's why the framework includes periodic audits and mandatory checkpoints, instead of leaving sensitivity as a good initial intention that fades under deadline pressure.